

CS CM122 Discussion

Apr 14 2023

Overview for today


Part 1. Textbook chapter 9

Part 2. Project 1 overview

Chapter 9

Trie of a set of patterns

- Each edge corresponds to a character
- Each path to a leaf (end node) is a string in the set of patterns
- Each node has unique outgoing edges corresponding to different characters
- To search for a string in the trie, start from the root and traverse the edges corresponding to the characters in the string
- Question: what is the time complexity for searching a string of length K in a trie?



Trie of patterns: "ananas", "and", "antenna", "banana", "bandana", "nab", "nana", "pan" (textbook 9.3)


Suffix trie

The suffix trie of a string Text is the trie formed from all its suffixes.

For example, the suffix trie of “ATAC\$” is the trie with set of patterns: “\$”, “C\$”, “AC\$”, “TAC\$”, “ATAC\$”.

Each path to a leaf is a unique suffix

The leaves are labelled with the starting positions of the suffix

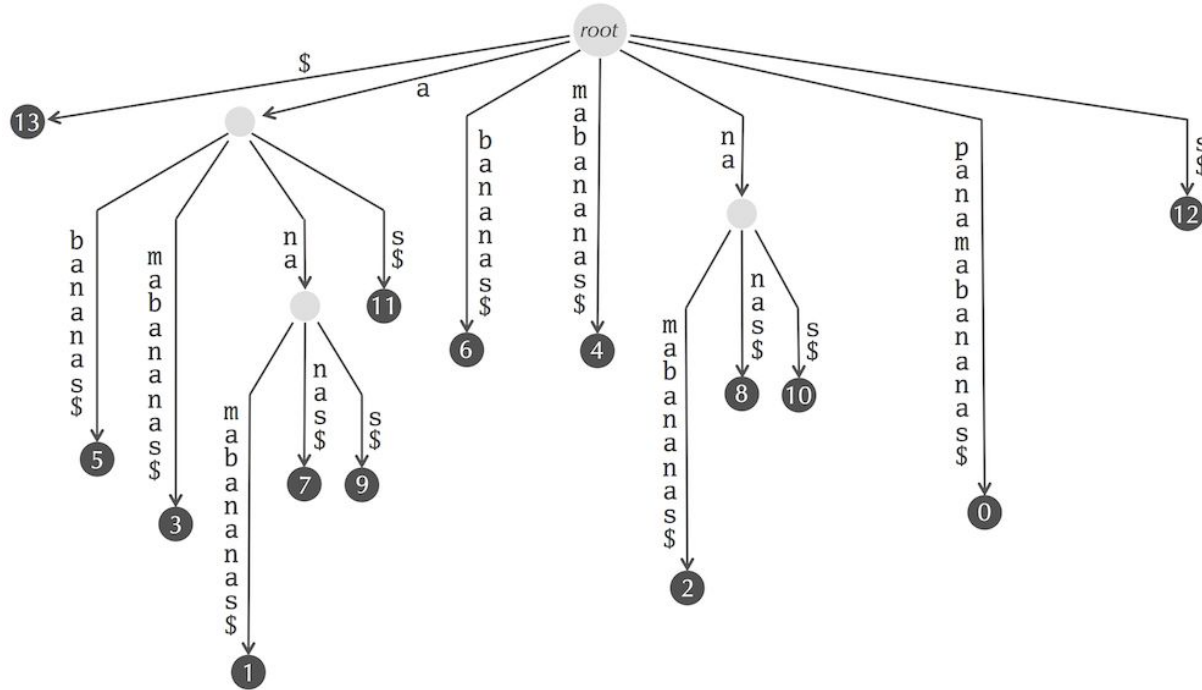


Suffix trie of text panamabananas\$ (textbook 9.4)

In-class whiteboard demonstration: why the “\$” sign is mandatory

- Suffix trie of string “PAPA”
- Suffix trie of string “PAPA\$”

Suffix tree





Merging single-branch paths in a suffix trie results in a suffix tree.

Each path to leaf still corresponds to a unique suffix. The number of leaves is therefore equal to the number of suffixes.

Suffix tree of text panamabananas\$ (textbook 9.5)

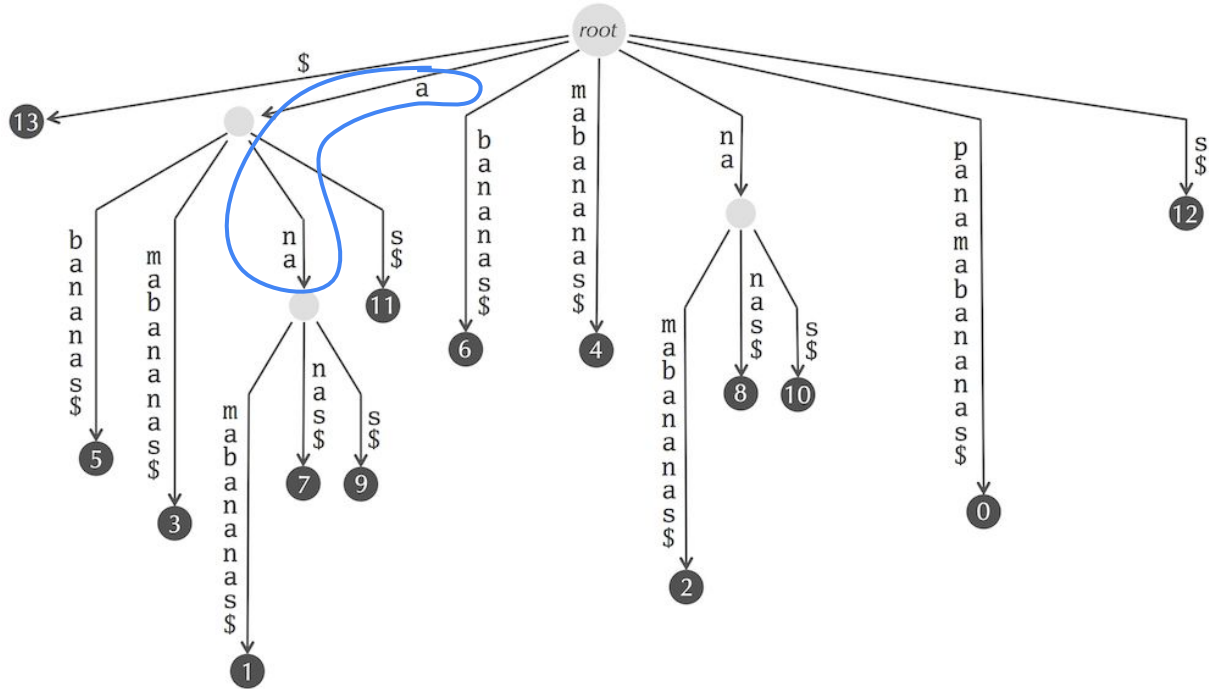
Converting a suffix trie to a suffix tree

- Traverse every node in the suffix trie while keeping track of its previous edge
- If you encounter a node with only 1 child, merge the two nodes and update the previous edge, then traverse the child(ren).
- If you encounter a node with multiple children, traverse every child instead
- Stop when you encounter a leaf



Longest repeat in a string

The longest repeat in a string is its longest suffix tree path that ends in a branch



Suffix array

All suffixes of a string are sorted. The resulting list of starting positions is the suffix array (no need to store the actual suffixes in memory)

Question: given a suffix array of length N, what is an efficient algorithm to search for a substring and what is its time complexity?

Starting Positions	Sorted Suffixes
13	\$
5	abananas\$
3	amabananas\$
1	anamabananas\$
7	ananas\$
9	anas\$
11	as\$
6	bananas\$
4	mabananas\$
2	namabananas\$
8	nanas\$
10	nas\$
0	panamabananas\$
12	s\$

Suffix array of text panamabananas\$ (textbook 9.6)

Burrows-Wheeler transform of “panamabananas\$”

Cyclic Rotations	$M(\text{"panamabananas\$"})$
panamabananas\$	\$ p a n a m a b a n a n a s
\$panamabananas	a b a n a n a s \$ p a n a m
s\$panamabanana	a m a b a n a n a s \$ p a n
as\$panamabanana	a n a m a b a n a n a s \$ p
nas\$panamabana	a n a n a s \$ p a n a m a b
anas\$panamaban	a n a s \$ p a n a m a b a n
nanas\$panamaba	a s \$ p a n a m a b a n a n
ananas\$panamab	b a n a n a s \$ p a n a m a
bananas\$panama	m a b a n a n a s \$ p a n a
abananas\$panam	n a m a b a n a n a s \$ p a
mabananas\$pana	n a n a s \$ p a n a m a b a
amabananas\$pan	n a s \$ p a n a m a b a n a
namabananas\$pa	p a n a m a b a n a n a s \$
anamabananas\$p	s \$ p a n a m a b a n a n a

In the original text, the first letter in a row follows the last letter in the same row.

Question: to generate the BWT do you need to store the entire matrix of cyclic rotations in the memory?

First-last property of a Burrows-Wheeler transform

p**a**₃**n****a**₂**m****a**₁**b****a**₄**n****a**₅**n****a**₆**s****\$**

The i^{th} occurrence of a letter in the last column is exactly the same instance in the original text as the i^{th} occurrence of the letter in the first column.

\$	p	a	n	a	m	a	b	a	n	a	n	a	s
a ₁	b	a	n	a	n	a	s	\$	p	a	n	a	m
a ₂	m	a	b	a	n	a	n	a	s	\$	p	a	n
a ₃	n	a	m	a	b	a	n	a	n	a	s	\$	p
a ₄	n	a	n	a	s	\$	p	a	n	a	m	a	b
a ₅	n	a	s	\$	p	a	n	a	m	a	b	a	n
a ₆	s	\$	p	a	n	a	m	a	b	a	n	a	n
b	a	n	a	n	a	s	\$	p	a	n	a	m	a ₁
m	a	b	a	n	a	n	a	s	\$	p	a	n	a ₂
n	a	m	a	b	a	n	a	n	a	s	\$	p	a ₃
n	a	n	a	s	\$	p	a	n	a	m	a	b	a ₄
n	a	s	\$	p	a	n	a	m	a	b	a	n	a ₅
p	a	n	a	m	a	b	a	n	a	n	a	s	\$
s	\$	p	a	n	a	m	a	b	a	n	a	n	a ₆

In-class whiteboard demonstrations: BWT procedures

- Inverting a transformed text
- Finding occurrences of a substring

How to find the original positions of matched substrings?

Have an array containing original positions of each suffix

$M(\text{Text})$	$\text{SUFFIXARRAY}(\text{Text})$
\$ p a n a m a b a n a n a s	13
a b a n a n a s \$ p a n a m	5
a m a b a n a n a s \$ p a n	3
a n a m a b a n a n a s \$ p	1
a n a n a s \$ p a n a m a b	7
a n a s \$ p a n a m a b a n	9
a s \$ p a n a m a b a n a n	11
b a n a n a s \$ p a n a m a	6
m a b a n a n a s \$ p a n a	4
n a m a b a n a n a s \$ p a	2
n a n a s \$ p a n a m a b a	8
n a s \$ p a n a m a b a n a	10
p a n a m a b a n a n a s \$	0
s \$ p a n a m a b a n a n a	12

How do you know how many times a letter occurs in the last column?

Have an array counting the occurrence of each letter in the last column up to each row

<i>i</i>	<i>FirstColumn</i>	<i>LastColumn</i>	LASTTOFIRST(i)	COUNT						
				\$	a	b	m	n	p	s
0	\$ ₁	s ₁	13	0	0	0	0	0	0	0
1	a ₁	m ₁	8	0	0	0	0	0	0	1
2	a ₂	n ₁	9	0	0	0	1	0	0	1
3	a ₃	p ₁	12	0	0	0	1	1	0	1
4	a ₄	b ₁	7	0	0	0	1	1	1	1
5	a ₅	n ₂	10	0	0	1	1	1	1	1
6	a ₆	n ₃	11	0	0	1	1	2	1	1
7	b ₁	a ₁	1	0	0	1	1	3	1	1
8	m ₁	a ₂	2	0	1	1	1	3	1	1
9	n ₁	a ₃	3	0	2	1	1	3	1	1
10	n ₂	a ₄	4	0	3	1	1	3	1	1
11	n ₃	a ₅	5	0	4	1	1	3	1	1
12	p ₁	\$ ₁	0	0	5	1	1	3	1	1
13	s ₁	a ₆	6	1	5	1	1	3	1	1
				1	6	1	1	3	1	1

Project 1

Format of data files

```
>read_0/1
ACGTATTTGAACTCCGGTATCTACACATTACGAGACGCATTATCAGCGTA
>read_0/2
TTAGGAAAGTCAAATGCATGGATCAGGGGCAAGATGCAGACACGGCTTAC
>read_1/1
AATGCGACGAAAGGCAGTACTTGTGCTCCATCTAGTTTGACGTATCCCAA
>read_1/2
AAGACGTTCTTGCTCAGCTTGAGAGCCCTCCGCGCCACGCAGTCACCCA
>read_2/1
ATGAGATACAGTACATGAGTGCTCCTCTACTGACACGTTTCGCTTTGCTC
>read_2/2
AGCCCTCCGCGCCACGCAGTCACCCAGTGCCGCTGATGCCAAGCACAG
>read_3/1
TTTAACGCGCTCTCCTTCCCGCTTCAGGGAATAATAGCAAGCGTGTTTT
```

Reads file: each line is a read,
each read may contain errors

```
>genome length:1000 generated:04132023_1331
TTATAAACACAGGACGAGCGCTCCGGATCAAAAACAACAGTCTGGCTAAACGAGTAACTCGACCCCGAGTGTGAGCAAT
CGTAGACGTCTGTGGTATTGGGCAAGGTTTTAGAAATTGCTATGGGCCCTATAGTCATTTGGGGCTTGCCTCCTATAGTT
CTCCGTATCCAGTTGTGCTAATGGGAGGTCGCCAGGCGGGGGACCAACTATGCCCCACAGGACAAATCTGACGCCGTGAT
TGCAGCCCAAGGTTTAAACGTAACGCGGCCCGCTTAATTTGGATATGTCGGTGGGTTCGGCATATGTAGATGCTT
GTTGTAACCGAGATGCCTCAGGCAGATACCTTAATGCGACGAAAGGCAGCACTTGTGCTCCATCTAGTTTGACGTATCCC
AAGGATGAGATACATACATGAGTGCTCCTCTACTGACACGTTTCGCTTTGCTACAGCAAAACATTAATCCAACGCAGTC
CGCAGGTATGGTGACTAGCGCAAAGTTTGTCTGTATCTTAGTAAGCCGTTAGTTTCGAAGACTGCCGCTACTCTGTTGAA
```

Genome file: a genome split
into several lines

```
>S129 C A
>S141 G A
>S219 G T
>S298 G T
>S369 C T
>S455 A G
>S460 A G
>S718 C A
>S863 T G
>I413 G
>I624 G
>D281 T
>D544 C
```

Your output should look like
this. S for substitutions, I for
insertions and D for deletions

Types of genetic mutations used for grading

https://commons.wikimedia.org/wiki/File:Deletion_Insertion_Substitution-en.svg

... AGCGTCGATGGAGATT ...



... AGCGT - - - - AGATT ...
... AGCGTAGATT ...

Deletion

... AGCGTCGATGGAGATT ...
... AGCGTCGACCATTTGGAGATT ...

Insertion

... AGCGTCGATGGAGATT ...
... AGCGTCGCTGGAGATT ...

Substitution

What data structure should you use to represent the genome?

Lecture 3

Index for L/3 (is BIG!)

- Intuition: Create an index (or phone book) for the genome.
- We can look up an entry quickly.

If $L=30$, each entry will have a key of length 10. Each entry will contain on average $N/4^{10}$ positions. (Approximately 3,000).

If $L=45$, each entry will have a key of length 15. Each entry will contain on average 3 positions.

Sequence	Positions
AAAAAAAAAA	32453, 64543, 76335
AAAAAAAAAC	64534, 84323, 96536
AAAAAAAAAG	12352, 32534, 56346
AAAAAAAAAT	23245, 54333, 75464
AAAAAAAAACA	
AAAAAAAACC	43523, 67543
...	
CAAAAAAAAA	32345, 65442
CAAAAAAAC	34653, 67323, 76354
...	
TCGACATGAG	54234, 67344, 75423
TCGACATGAT	11213, 22323
...	
TTTTTTTTTG	64252
TTTTTTTTTT	64246, 77355, 78453

You can create a hash table containing positions of all unique k-mers in the genome.

You can also use tries or BWT.

How do you map a read to the genome?


Assume you have a read of length L and a hash table k -mer positions.

You can cut the read into L/k fragments and try to find the position of each fragment. x

Genome



Read (with 1 substitution in the middle)




Mapped to a position A

Mapped to a position $A + k$

Not found in hash or
mapped to somewhere else

How do you identify mutations from aligned reads?

```
Ref:  TTTTGTCACCCTCAACAACGACGCT
Read:  .....
Read:  ....TCACCCTCAACAAACGACGCT
Read:  .....AAACGACGCT
Read:  ....ACCCTCAACAAACGACGCT
Read:  TTTTGTCACCCTCAACAAACGACGCT
Read:  ....CCCTCAACAAACGACGCT
Read:  .....
Read:  ....CCTCAACAAACGACGCT
```



After aligning reads to the reference genome, your program should be able to compare the reference and the reads at each position and check if a different nucleotide is present in the reads.

Advanced read mapping techniques

Even better fragment-based alignment

Published: 04 March 2012

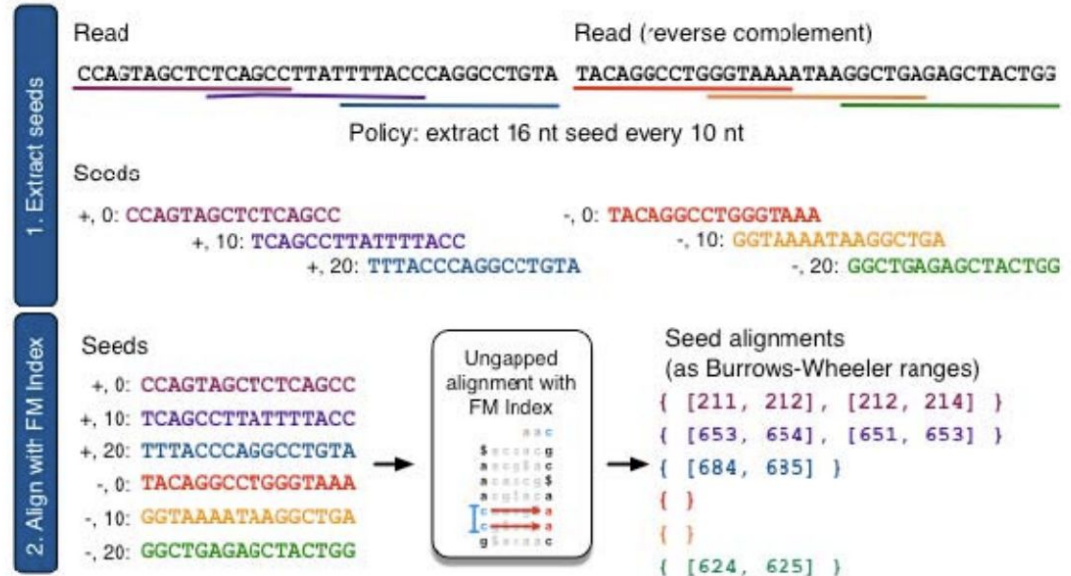
Fast gapped-read alignment with Bowtie 2

Ben Langmead  & Steven L. Salzberg

Nature Methods 9, 357–359 (2012) | [Cite this article](#)

82k Accesses | 27158 Citations | 129 Altmetric | [Metrics](#)

Instead of cutting a read of length L into L/k fragments, apply a sliding window of length k across the read (this allows you to get more than L/k fragments)



What if your fragments are mapping to many different positions?

JOURNAL ARTICLE

The Subread aligner: fast, accurate and scalable read mapping by seed-and-vote

Yang Liao, Gordon K. Smyth, Wei Shi

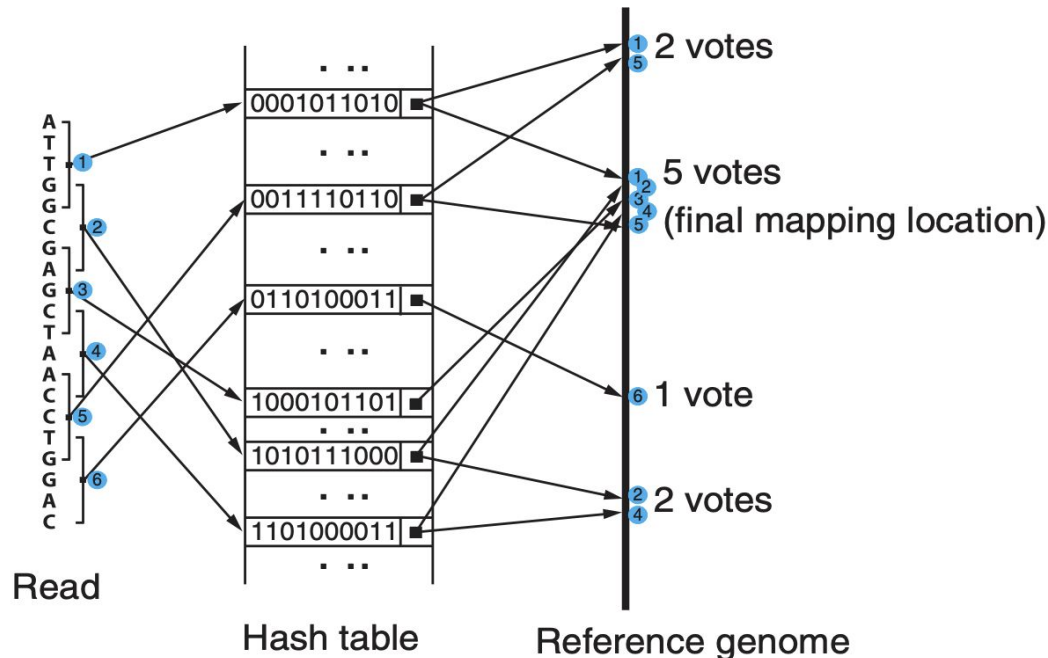
Nucleic Acids Research, Volume 41, Issue 10, 1 May 2013, Page e108, <https://doi.org/10.1093/nar/gkt214>

Published: 03 April 2013 [Article history](#)

Each fragment in a read casts a “vote” on each of the potential genomic regions where it came from.

The region with the majority of the votes from all fragments wins.

B



Needleman-Wunsch alignment for two sequences (in Chapter 5 of textbook)

This allows you to find any arbitrary indels.

Needleman-Wunsch

match = 1 mismatch = -1 gap = -1

		G	C	A	T	G	C	G
	0	-1	-2	-3	-4	-5	-6	-7
G	-1	1	0	-1	-2	-3	-4	-5
A	-2	0	0	1	0	-1	-2	-3
T	-3	-1	-1	0	2	1	0	-1
T	-4	-2	-2	-1	1	1	0	-1
A	-5	-3	-3	-1	0	0	0	-1
C	-6	-4	-2	-2	-1	-1	1	0
A	-7	-5	-3	-1	-2	-2	0	0

How would indels look like on your read?

Genome




Read



Deletion: some fragments of your read will map to positions that are more far apart on the genome

Genome



Read



Insertion: some fragments of your read will map to positions that are closer on the genome