



Algorithms in Bioinformatics

Spring 2023

Lecture 2

Eleazar Eskin

University of California, Los Angeles

Administrative updates and reminders

- Check course website for latest announcements:
<https://bruinlearn.ucla.edu/courses/160773>

will also post lecture slides before class

- Discussion sections will be primarily covering chapter 1 this week
 - 1A – Fri 12-1:50pm BROAD 2100A
 - 1B – Fri 12-1:50pm KAPLAN A65
 - 1C – Fri 2-3:50pm DODD 175
 - 1D – Fri 2-3:50pm KAPLAN 169
- Everyone on waitlist at the end of class Tue should have a PTE in MyUCLA
- If still need a PTE, email a TA or instructor with UID and if not on waitlist preferred discussion section(s)

Textbook

- Accounts have been created for receipts uploaded by yesterday
- Use same email with online book site as given in the receipt upload form
- Additional receipts for print book (<https://www.bioinformaticsalgorithms.org/>) uploaded to <https://forms.gle/iA7JLc7LK2a3Qiuv8> by Tue April 11th will be processed for access to online site by Wed. Apr 12th
- Link for online only purchase <https://stepik.org/a/170824>
- Form to transfer previous purchase of textbook to our course instance for free:
<https://forms.gle/ms6BRrC6o5eevBGf7>



Upcoming due dates

- HW1 due Thur. 4/13 (recommend to finish earlier)
- Post Question on Paper 1 due Thur. 4/13




Sequencing + Read Mapping

Lecture 2.

April 6th, 2023

Active Research Problem: Short Read Re-sequencing Where are my mutations?



Sequencing Technology



Illumina / Solexa
Genetic Analyzer 1G
1000 Mb/run, 35bp reads

```
AGAGCAGTCGACAGGTA
TAGTCTACATGAGATCG
ACATGAGATCGGTAGAG
CCGTGAGATCGACATGA
TAGCCAGAGCAGTCGAC
AGGTATAGTCTACATGA
GATCGACATGAGATCGG
TAGAGCCGTGAGATCGA
CATGATAGCCAGAGCAG
TCGACAGGTATAGTCTA
CATGAGATCGACATGAG
ATCGGTAGAGCCGTGAG
ATCGACATGATAGCCAG
AGCAGTCGACAGGTATA
GTCTACATGAGATCGAC
ATGAGATCGGTAGAGCC
GTGAGATCGACATGATA
GCCAGAGCAGTCGACAG
GTATAGTCTACATGAGA
TCGACATGAGATCGGTA
GAGCCGTGAGATCGACA
TGATAGCCAGAGCAGTC
GACAGGTATAGTCTACA
TGAGATCGACATGAGAT
CGGTAGAGCCGTGAGAT
CGACATGATAGCCAGAG
CAGTCGACAGGTATAGT
CTACATGAGATCGACAT
GAGATCGGTAGAGCCGT
GAGATCGACATGATAGC
```


■ Next generation sequencing.

- Cheap sequencing.
- “Short Reads”

Short Read Sequencing Problem (A Computer Science Problem)

Full DNA Sequence

AGAGC**A**GTTCGACAG**G**GTA
TAG**T**CTACATGAGATC**G**
ACATGAGATC**G**GTAGAG
CCGTGAGATC**G**ACATGA
TAG**C**CAGAGC**A**GTTCGAC
A**G**GTATAG**T**CTACATGA
GATC**G**ACATGAGATC**G**G
TAGAGC**C**GTGAGATC**G**A
CATGATAG**C**CAGAGC**A**G
TCGACA**G**GTATAG**T**CTA
CATGAGATC**G**ACATGAG
ATC**G**GTAGAGC**C**GTGAG
ATC**G**ACATGATAG**C**CAG
AGC**A**GTTCGACAG**G**GTATA
G**T**CTACATGAGATC**G**AC
ATGAGATC**G**GTAGAGC**C**
GTGAGATC**G**ACATGATA
G**C**CAGAGC**A**GTTCGACAG**G**
GTATAG**T**CTACATGAGA
TC**G**ACATGAGATC**G**GTA
GAGC**C**GTGAGATC**G**ACA
TGATAG**C**CAGAGC**A**GTTC
GACA**G**GTATAG**T**CTACA
TGAGATC**G**ACATGAGAT
C**G**GTAGAGC**C**GTGAGAT
C**G**ACATGATAG**C**CAGAG
C**A**GTTCGACAG**G**GTATAG**T**
CTACATGAGATC**G**ACAT
GAGATC**G**GTAGAGC**C**GT
GAGATC**G**ACATGATAG**C**



- Short read sequencers generate random short substrings from the DNA sequence of a certain length.

ATGAGATCGGTAGAGCCGTGAGAT
GAGCAGTCGACAGGTATAGTCTAC
AGAGCAGTCGACAGGTATAGTCTA
TGAGATCGACATGATAGCCAGAGC
TAGCCAGAGCAGTCGACAGGTATA
GATAGCCAGAGCAGTCGACAGGTA
GAGATCGACATGATAGCCAGAGCA
GCAGTCGACAGGTATAGTCTACAT
AGCAGTCGACAGGTATAGTCTACA
TCGACATGAGATCGGTAGAGCCGT
CAGTCGACAGGTATAGTCTACATG
GAGATCGACATGATAGCCAGAGCA
GTAGAGCCGTGAGATCGACATGAT

How do we
recover the
original
sequence?

Short Reads Difficulties

ATGAGATCGGTAGAGCCGTGAGAT
GAGCAGTCGACAGGTATAGTCTAC
AGAGCAGTCGACAGGTATAGTCTA
TGAGATCGACATGATAGCCAGAGC
TAGCCAGAGCAGTCGACAGGTATA
GATAGCCAGAGCAGTCGACAGGTA
GAGATCGACATGATAGCCAGAGCA
GCAGTCGACAGGTATAGTCTACAT
AGCAGTCGACAGGTATAGTCTACA
TCGACATGAGATCGGTAGAGCCGT
CAGTCGACAGGTATAGTCTACATG
GAGATCGACATGATAGCCAGAGCA
GTAGAGCCGTGAGATCGACATGAT

- We don't know where each read comes from!
- Can't identify where the mutations are!
- What do we do?

Key Idea: “Re”-Sequencing

We know that my genome is very close to the Human genome.

My Genome:

TACATGAGATC**G**ACATGAGATC**G**GTAGAGCC**C**GTGAGATC

A Sequence Read:

TCGACATGAGATCGGTAGAGCCGT

The Human Genome:

TACATGAGATC**C**ACATGAGATC**T**GTAGAGCC**T**GTGAGATC
TC**C**GACATGAGATC**T**GGTAGAGCC**T**

Recovered Sequence:

TACATGAGATC**G**ACATGAGATC**G**GTAGAGCC**C**GTGAGATC

“Re”-Sequencing Challenges (Why do we need Computer Science?)

- Sequences are long!
 - **Human Genome is 3,000,000,000 long.**
- Sequencers generate many reads!
 - **A single run generates over 300,000,000 reads.**
- We need efficient algorithms to “map” each read to its location in the genome.
 - **A trivial mapping algorithm will take thousands of years to compute for a genome.**

There are other challenges which we are not mentioning.

Trivial Mapping Algorithm

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

TCGACATGAGATCGGTAGAGCCGT

- We can slide our read along the genome and count the total mismatches between the read and the genome.
- If the mismatches are below a threshold, we say that it is a match.

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
TCGACATGAGATCGGTAGAGCCGT
↑↑ ↑↑↑ ↑↑↑↑↑ ↑↑↑↑↑ ↑↑

Total of 18 mismatches. Not below threshold. Not a match.

Trivial Mapping Algorithm

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

TCGACATGAGATCGGTTAGAGCCGT

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
TCGACATGAGATCGGTTAGAGCCGT



Total of 15 mismatches. Not below threshold. Not a match.

Trivial Mapping Algorithm

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

TCGACATGAGATCGGTTAGAGCCGT

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
TCGACATGAGATCGGTTAGAGCCGT



Total of 23 mismatches. Not below threshold. Not a match.

Trivial Mapping Algorithm

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

TCGACATGAGATCGGTTAGAGCCGT

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
TCGACATGAGATCGGTTAGAGCCGT



Total of 23 mismatches. Not below threshold. Not a match.

Trivial Mapping Algorithm

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

TCGACATGAGATCGGTAAGAGCCGT

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
TCGACATGAGATCGGTAAGAGCCGT



Total of 3 mismatches. Below threshold. A match!

Complexity of Trivial Algorithm

- 3,000,000,000 length genome (N)
- 300,000,000 reads to map (M)
- Reads are of length 30 (L)
- Number of mismatches allowed is 2 (D).
- Each comparison of match vs. mismatch takes 1/1,000,000 seconds (t).

Total Time = $N * M * L * t = 27,000,000,000,000$ seconds or 854,164 years!




Some observations

- Most positions in the genome match very poorly.
- We are looking for only a few mismatches.
(D is small)
- A substring of our read will match perfectly.

Perfect Matching Read Substrings

Three “worst” possible cases for placement of mutations.



- In each case, there is a perfect match of $L/3$.

Finding a perfect match of length $L/3$

- Intuition: Create an index (or phone book) for the genome.
- We can look up an entry quickly.

If $L=30$, each entry will have a key of length 10. Each entry will contain on average $N/4^{10}$ positions. (Approximately 3,000).

If $L=45$, each entry will have a key of length 15. Each entry will contain on average 3 positions.

Sequence	Positions
AAAAAAAAAA	32453, 64543, 76335
AAAAAAAAAC	64534, 84323, 96536
AAAAAAAAAG	12352, 32534, 56346
AAAAAAAAAT	23245, 54333, 75464
AAAAAAAAACA	
AAAAAAAAACC	43523, 67543
...	
CAAAAAAAAA	32345, 65442
CAAAAAAAAAAC	34653, 67323, 76354
...	
TCGACATGAG	54234, 67344, 75423
TCGACATGAT	11213, 22323
...	
TTTTTTTTTTG	64252
TTTTTTTTTTT	64246, 77355, 78453

Complexity of Indexing Algorithm

- We need to look up each third of the read in the index.
- For $L=30$, our index will contain entries of length 10. Each entry will contain on average $N/(4^{L/3})$ or 3,000 positions.
- For each position, we need to compute the number of mismatches.
- Our running time is $L * M * 3 * N / (4^{L/3}) * T = 81,000,000$ seconds or 937 days.
- If $L=45$, then the time is 81,000 seconds or 22.5 hours.

More problems: Sequencing Errors

- Each sequence read can have some random errors.

My Genome:

TACATGAGATC**G**ACATGAGATC**G**GTAGAGC**C**GTGAGATC

A Sequence Read:

TCGACATGAGATCGGTAGAA**A**CCGT

The Human Genome:

TACATGAGATC**[]**ACATGAGATC**[]**GTAGAG**[]****[]**GTGAGATC
TCGACATGAGATCGGTAGAA**[]**CCGT

Recovered Sequence:

TACATGAGATC**G**ACATGAGATC**G**GTAGA**A****C**GTGAGATC

Sequencing Errors: Solution

- Collect redundant data.

My Genome:

TACATGAGATC**G**ACATGAGATC**G**GTAGAGC**C**GTGAGATC

Sequence Reads:

TCGACATGAGATC**G**GTAGAG**A**CCGT

GACA**A**GAGATC**G**GTAGAGCCGTGA

TGAGATC**G**AGAGCCGTGAGATC

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

TC**G**ACATGAGATC**G**GTAGAG**A****C**GT

GACA**A**GAGATC**G**GTAGAGC**C**GTGA

TGAGATC**G****G**AGAGC**C**GTGAGATC

Recovered Sequence:

TACATGAGATC**G**ACATGAGATC**G**GTAGAGC**C**GTGAGATC


How much coverage do we need?

- If error rate is e , and we are going to predict the consensus sequence, what is the error rate if the coverage is 3.
- We will make a prediction with an error if two out of our three reads have an error in the same place.
- This is approximately $3e^2$.

“Re”-Sequencing Problems

The Human Genome:

TACATGAGATCCACATGAGATCTGTACATGAGATCCACAT



Repeated Region

My Genome:

TACATGAGATC**G**ACATGAGATCGGTACATGAGATCCACAT

A Sequence Read:

ACATGAGATC**G**ACAT

The Human Genome:

TACATGAGATC**G**ACATGAGATCTGTACATGAGATC**G**ACAT
ACATGAGATC**G**ACAT ACATGAGATC**G**ACAT

Error!

Recovered Sequence:

TACATGAGATC**G**ACATGAGATCGGTACATGAGATC**G**ACAT



“Re”-Sequencing Problems

The Human Genome:

TACATGAGATCCACATGAGATCTGTACATGAGATCCACAT

My Genome:

TACATGAG**GGGGGGGG**GAGATCGGTACATGAGATCCACAT

A Sequence Read:

GAG**GGGGGGGG**

The Human Genome:

TACATGAGATCCACATGAGATCTGTACATGAGATCCACAT
GAG**GGGGGGGG**

Too many mismatches to match the read to the reference.
Since we don't know where it came from, we can't identify
the difference in the target sequence.

“Re”-Sequencing: Insertions

My Genome:

TACATGAGATCCACATAGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

CCACATAGAGATCTGTAGAGCTGT

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT



TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT



How do we deal with this case?

“Re”-Sequencing: Insertions

My Genome:

TACATGAGATCCACATAGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

CCACATAGAGATCTGTAGAGCTGT

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT



TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT



Solution: Add Insertion to the Human Genome

TACATGAGATCCACAT-GAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT



Many other challenges

- Coverage of sequence reads is not uniform
 - **Some places we have many reads, while some we have fewer. How do we design an approach so we can always recover the sequence.**
- Large memory requirements
 - **We need to fit our index into RAM. Often tens of Gigabytes or greater.**



Sequencing Coverage

Lecture 2.

April 6th, 2023

(Slides from Jae-Hoon Sul)



Sequence Mapping Coverage

- If a genome is length N (human is 3,000,000,000), and the total length of all sequence reads collected is M , the coverage ratio is defined at M/N .
- Often written with an “x”. For example, 10x or 20x coverage.



Sequencing Coverage Statistics

- If length of the genome is N the probability of the event that a single read position starts at a single position in the genome is $1/N$ (very small).
- If the number of reads is K , the total number of read positions that start at a single genome position is the number of times that an event with probability $1/N$ happens out of K trials.
- Poisson distribution.



Sequencing Coverage Statistics


- If length of the genome is N the probability of the event that a single read of length L position spanning a single position in the genome approximately L/N (also very small).
- If the sum of the length of all K reads of length L is $M=K*L$, the total number of read positions that span a single genome position is approximately the number of times that an event with probability $1/N$ happens out of M trials.
- Approximately Poisson distribution.

Poisson Distribution

- Discrete probability distribution to compute probability of (rare) events given known mean
- Only one parameter: λ , mean of distribution
- Probability Mass Function

$$\Pr(N_t = k) = \frac{e^{-\lambda} \lambda^k}{k!}$$

- Mean = λ
- Variance = λ



Poisson Distribution to Sequencing Coverage

- $\lambda = M/N$.
- Probability that exactly X reads span a certain position.
 - **dpois(X, λ)**
- Probability that X or fewer reads span a certain position.
 - **ppois(X, λ)**
- At least $Y\%$ of the genome is covered with this much or fewer reads
 - **qpois(Y, λ)**

Poisson and Sequencing Coverage

- Probability that X or fewer reads span a certain position.

- $\mathbf{ppois}(X, \lambda) = \sum_{i=0}^X \mathbf{dpois}(i, \lambda)$

Coverage examples

- For human genome ($L=3,000,000,000$) sequenced at 30x coverage, what is the probability that a specific location has exactly 30 coverage?
- **$\lambda=30$ $\text{dpois}(30,\lambda)=\text{dpois}(30,30)=0.072$**
- What is the probability that a specific location has at least 30 coverage?
- **$1-\text{ppois}(29,\lambda)=1-\text{ppois}(29,30)=0.524$**
- What is the probability that a specific location has at least 10 coverage?
- **$1-\text{ppois}(9,30)=0.9999929$**

Coverage examples

- For human genome ($L=3,000,000,000$) sequenced at 30x coverage, what is the probability that a specific location has exactly one read spanning it?
- **$\lambda=30$ $\text{dpois}(1,\lambda)=2.9 \times 10^{-12}$**
- What is the probability that a specific location has at least 6 coverage?
- **$\lambda=30$ $1-\text{ppois}(5,\lambda)=.99999$**
- How many positions in the genome have less than 6 coverage ?
- **$3,000,000,000 * \text{ppois}(5,\lambda)=67.7$**

Diploid Coverage

- Since humans have 2 chromosomes each read comes from one chromosome at random. If a position in the reference is covered by Y reads, the probability that X of the reads come from the first chromosome follows the binomial distribution with parameter .5.

□ **`dbinom(X,Y,0.5)`**

- At least X coverage for each chromosome out of Y reads

$$\sum_{i=X}^{Y-X} \text{dbinom}(i, Y, 0.5)$$


Binomial Distribution

- Discrete probability distribution to compute probability of having X successes in Y trials
- Example: What's the probability of having k heads in n tosses with fair coin ($p = 0.5$)?

- Probability Mass Function

$$\binom{n}{k} p^k (1-p)^{n-k}$$

- Mean = $n \cdot p$
- Variance = $n \cdot p \cdot (1-p)$



Diploid Coverage Examples

- If a position is covered by 10 reads, what is the probability that exactly 3 reads come from the first chromosome?
- **$\text{dbinom}(3,10,.5)=.117$**
- If a position is covered by 10 reads, what is the probability that at least 4 reads come from the first chromosome?
- **$1-\text{pbinom}(3,10,.5)=.828$**
- If a position is covered by 10 reads, what is the probability that at least 4 reads come from each chromosome?
- **$\text{dbinom}(4,10,.5)+\text{dbinom}(5,10,.5)+\text{dbinom}(6,10,.5)=.656$**

Minimum Diploid Coverage

- If we want the sequence coverage is $\lambda=M/N$, the portion of the genome that has at least X coverage of each chromosome is

$$\sum_{i=2X}^{\infty} \text{dpois}(i, \lambda) \sum_{j=X}^{i-X} \text{dbinom}(j, i, 0.5)$$

Diploid Coverage Examples

- If genome is covered with coverage 30, what is the probability that a position will have at least 10 reads from each chromosome?

$$\sum_{i=20}^{\infty} \text{dpois}(i, 30) \sum_{j=10}^{i-10} \text{dbinom}(j, i, 0.5)$$



SNP Calling

- Inferring single base differences from sequencing.
- Several challenges:
 - **Sequencing errors**
 - **Alignment “mapping” problems**
 - **Statistical Uncertainty**



SNP Calling Standard Approaches

■ Consensus Algorithm

- **Map reads to genome**
- **Place read in best mapping position (randomly break ties)**
- **SNP call is based on majority vote.**

■ Probabilistic Algorithm

- **Map reads to genome**
- **Place read in best mapping position (randomly break ties)**
- **Compute “posterior probability”**

■ Mapping uncertainty methods

- **Map reads to genome**
- **Record mapping uncertainty**
- **Compute “posterior probability” incorporating mapping uncertainty**

Sequencing Errors

- Each sequence read can have some random errors.

My Genome:

TACATGAGATC**G**ACATGAGATC**G**GTAGAGCC**C**GTGAGATC

A Sequence Read:

TCGACATGAGATC**A**CGTAGA**A**CCGT

The Human Genome:

TACATGAGATC**C**ACATGAGATC**C**GTAGAG**C****C**GTGAGATC
TCGACATGAGATC**C**CGTAGA**A**CCGT

Recovered Sequence:

TACATGAGATC**G**ACATGAGATC**G**GTAGA**A****C**CGTGAGATC

Consensus Algorithm

- Take majority vote.

My Genome:

TACATGAGATC**G**ACATGAGATC**G**GTAGAGC**C**GTGAGATC

Sequence Reads:

TCGACATGAGATC**G**GTAGA**A**CCGT

GACA**A**GAGATC**G**GTAGAGCCGTGA

TGAGATC**G**G**A**GAGCCGTGAGATC

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

TC**G**ACATGAGATC**G**GTAGA**A****C**GT

GACA**A**GAGATC**G**GTAGAGC**C**GTGA

TGAGATC**G**G**A**GAGCC**C**GTGAGATC

Recovered Sequence:

TACATGAGATC**G**ACATGAGATC**G**GTAGAGC**C**GTGAGATC

How much coverage do we need?

- If error rate is e , and we are going to predict the consensus sequence, what is the error rate if the coverage is X .
- We will make a prediction with an error more than $X/2 - 1$ out of the X reads have an error in the same place.
- **$1 - \text{pbinom}(X/2, X, e)$**

How much coverage do we need?

- If error rate is e , and we are going to predict the consensus sequence, what is the error rate if the coverage is 3.
- We will make a prediction with an error if two out of our three reads have an error in the same place.

$$\text{pbinom}(2, 3, e) = e^3 + \binom{3}{2}(1-e)e^2$$

- This is approximately $3e^2$.



Diploid Sequencing

- Humans have 2 chromosomes.
- Each chromosome may have a different SNP.
- Some reads come from 1 chromosome, some come from other chromosome.
- Why does consensus method not work?
- How do we address this problem?



Insertions and Deletions

Lecture 2.

April 6th, 2023

“Re”-Sequencing: Insertions

My Genome:

TACATGAGATCCACATAGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

CCACATAGAGATCTGTAGAGCTGT

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT



TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT



How do we deal with this case?

“Re”-Sequencing: Insertions

My Genome:

TACATGAGATCCACATAGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

CCACATAGAGATCTGTAGAGCTGT

The Human Genome:

TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT



TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT




Solution: Add Insertion to the Human Genome

TACATGAGATCCACAT-GAGATCTGTAGAGCTGTGAGATC
CCACATAGAGATCTGTAGAGCTGT

Indel in Middle of Read

If indel is in the middle of read.



- Both outside regions of size $L/3$ will match perfectly.
- Because of coverage, indel will be in middle at least for one read.
- Important: Middle distance will be $L/3+1$ or $L/3-1$

“Re”-Sequencing: Insertions

My Genome:

TACATGAGATCCACATAGAGATCTGTAGAGCTGTGAGATC

A Sequence Read:

CCACATAGAGATCTGTAGAGCTGT

The Human Genome:


TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

CCACATAGAGATCTGTAGAGCTGT



TACATGAGATCCACATGAGATCTGTAGAGCTGTGAGATC

CCACATAGAGATCTGTAGAGCTGT



Insertion Point

Indel Algorithms

■ Trivial Algorithm

- **Try all insertion points for a read**
- **If read matches (with insertion) below number of mismatches, then we declare a match and identify and indel**

■ More Efficient Algorithm

- **Look for perfect match in first part of read**
- **Try insertion point at point of first mismatch**
- **More complicated but faster**

■ More accurate Algorithm

- **Perform alignment between read and reference**

■ Extremely Accurate Algorithm

- **Align all reads with indel together.**
- **Multiple Sequence Alignment!**



“Re”-Sequencing + Burroughs Wheeler Transform

Lecture 2.

April 6th, 2023

(Some slides from Ben Langmead)

Index for L/3 (is BIG!)

- Intuition: Create an index (or phone book) for the genome.
- We can look up an entry quickly.

If $L=30$, each entry will have a key of length 10. Each entry will contain on average $N/4^{10}$ positions. (Approximately 3,000).

Sequence	Positions
AAAAAAAAAA	32453, 64543, 76335
AAAAAAAAAC	64534, 84323, 96536
AAAAAAAAAG	12352, 32534, 56346
AAAAAAAAAT	23245, 54333, 75464
AAAAAAAAACA	
AAAAAAAAACC	43523, 67543
...	
CAAAAAAAAA	32345, 65442
CAAAAAAAC	34653, 67323, 76354
...	
TCGACATGAG	54234, 67344, 75423
TCGACATGAT	11213, 22323
...	
TTTTTTTTTG	64252
TTTTTTTTTT	64246, 77355, 78453

If $L=45$, each entry will have a key of length 15. Each entry will contain on average 3 positions.




Indexing a genome

- To find exactly matching substrings, we need to build an index for the whole genome.
- Problem: The genome is BIG!

Indexing


- Genome indices can be big. For human:



> 35 GBs

6	\$
5	A\$
3	ANA\$
1	ANANA\$
0	BANANA\$
4	NA\$
2	NANA\$

> 12 GBs



> 12 GBs

- Large memory requirement implications
 - **Requires large memory machine (expensive)**
 - **Partition genome and index each part (slow)**



Memory Efficient but Slow Algorithm

- Store just the sequence
 - **4 DNA bases per byte (2 bits each)**
 - **3,000,000,000 / 4 ~ 750 MB**
- When looking up string, just loop through the sequence.
- Very slow, but very memory efficient!

Burrows-Wheeler Transform

- http://en.wikipedia.org/wiki/Burrows-Wheeler_transform
- Reversible permutation used originally in compression




- Once $BWT(T)$ is built, *all else shown here is discarded*

□ **Matrix will be shown for illustration only**


Burrows-Wheeler Transform

- Store only last column
- First column can be recovered by counting symbols in last column because it is sorted



Burrows-Wheeler Transform


- Property that makes $BWT(T)$ reversible is “LF Mapping”
 - i^{th} occurrence of a character in Last column is same *text* occurrence as the i^{th} occurrence in First column



Burrows-Wheeler Transform

- Property that makes $BWT(T)$ reversible is “LF Mapping”

- i^{th} occurrence of a character in Last column is same *text* occurrence as the i^{th} occurrence in First column




Burrows-Wheeler Transform

- To recreate T from BWT(T), repeatedly apply rule:

$$\mathbf{T} = \mathbf{BWT}[\mathbf{LF}(i)] + \mathbf{T}; i = \mathbf{LF}(i)$$

- Where **LF**(i) maps row i to row whose first character corresponds to i's last per LF Mapping



- Could be called “unpermute” or “walk-left” algorithm


FM Index

- Ferragina & Manzini propose “FM Index” based on BWT
- Observed:
 - **LF Mapping also allows *exact matching* within T**
 - **LF(i) can be made fast with *checkpointing***
 - **...and more (see FOCS paper)**
- Ferragina P, Manzini G: Opportunistic data structures with applications. *FOCS. IEEE Computer Society; 2000.*
- Ferragina P, Manzini G: An experimental study of an opportunistic index. *SIAM symposium on Discrete algorithms*. Washington, D.C.; 2001.

Exact Matching with FM Index

- Look up pattern in reverse.
- Use 2 pointers to represent range of matches.
- Find first valid match for next symbol in range.


□ **Example: searching for “caa”**



Exact Matching with FM Index

- Look up pattern in reverse.
- Use 2 pointers to represent range of matches.
- Find first valid match for next symbol in range.


□ **Example: searching for “aac”**



Exact Matching with FM Index

■ If no match...

□ **Example: searching for “gac”**




■ Pointers will get lost.

■ FM index can quickly check for a match.

Where in sequence is the match?

- Use “walk-left” to build sequence to start
- Count number of sequences

□ **Example: searching for “aac”**



- Number of moves back is start position of match
- **Example: “aac” is in position 2.**

Where in sequence is match?

- “walk-left” to start of sequence is slow
- Requires on average $N/2$ steps to reach start.
- Alternate strategy: keep index of positions.
 - **Example: searching for “aac”**


	\$	a	c	a	a	c	g	6
→	a	a	c	g	\$	a	c	2
	a	c	a	a	c	g	\$	0
	a	c	g	\$	a	c	a	3
	c	a	a	c	g	\$	a	1
	c	g	\$	a	c	a	a	4
	g	\$	a	c	a	a	c	5
	aac							

- Problem: requires as much storage as hashtable!

Where in sequence is match?

- Key Idea: Store fraction of array (sampling)
- Only store some positions and “walk-left”
- Combines two previous strategies

□ **Example: searching for “aac”**



Position =
number of moves
+ position in array


For "aac" = 1 + 1 = 2

- How many values to store provides defines time/space tradeoff.

“walk-left” optimization

- Each “walk-left” requires counting previous occurrences of symbol in BWT

□ **Example: searching for “aac”**



- Requires counting occurrences in N/2 length string
- Really slow!

“walk-left” optimization

- Idea: use checkpoints to store previous counts
 - Example: searching for “aac”

	\$	a	c	a	a	c	g		A:0	C:0	G:1	T:0
→	a	a	c	g	\$	a	c					
	a	c	a	a	c	g	\$					
	a	c	g	\$	a	c	a		A:1	C:1	G:1	T:0
	c	a	a	c	g	\$	a					
	c	g	\$	a	c	a	a					
	g	\$	a	c	a	a	c		A:3	C:2	G:1	T:0
	aac											

- Requires counting occurrences only until checkpoint.
- Really fast!

FM Index is Small (Bowtie)

■ Entire FM Index on DNA reference consists of:


- **BWT (same size as T)**
- **Checkpoints (~15% size of T)**
- **SA sample (~50% size of T)**

Assuming 2-bit-per-base encoding and no compression, as in Bowtie


Assuming a 16-byte checkpoint every 448 characters, as in Bowtie

Assuming Bowtie defaults for suffix-array sampling rate, etc


■ Total: ~1.65x the size of T




~1.65x



>45x



>15x



>15x



Reference Paper

- Langmead B, Trapnell C, Pop M, Salzberg SL. Ultrafast and memory-efficient alignment of short DNA sequences to the human genome. *Genome Biology* 10:R25.
 - (Some slides from paper)